

CSE440: Natural Language Processing II

Dr. Farig Sadeque
Associate Professor
Department of Computer Science and Engineering
BRAC University

Lecture 4: Word Representations

Outline

- Co-occurrence (SLP 6)
- TF-IDF (SLP 6)
- Embeddings (SLP 6 and lecture)

Intro

- Computers do not understand semantics
- Representation of text needs to include some sort of semantic information

Representation

- Sentence-level representation problems
- Co-occurrence
- TF-IDF
- Embeddings

Problems with BoW

- Too sparse
 - What's wrong with sparsity?
- Completely ignores word order
- Almost no semantic information preserved
- But, works pretty well!

More problem with sentence-level representation

- Dogs chew snacks
- Canines eat treats

More problem with sentence-level representation

- Dogs chew snacks
- Canines eat treats

<i>Documents</i>	<i>Features</i>					
	$f_{canines}$	f_{chew}	f_{dogs}	f_{eat}	f_{snacks}	f_{treats}
<i>dogs chew snacks</i>	0	1	1	0	1	0
<i>canines eat treats</i>	1	0	0	1	0	1

- No feature overlap whatsoever. If we try to calculate similarity, they are 100% dissimilar. But are they?
- Solution: instead of creating sentence level representations, let's go to smaller units i.e. words

A smarter sentence representation: TF-IDF

TF-IDF: Term Frequency - **Inverse** Document Frequency

Intuition: An informative term should:

- Occurs many times in some specific contexts (TF)
- Does not occur in every context (IDF)

Examples:

- high TF **vector** is informative; it's frequent in these slides
- high DF **the** is uninformative; it's frequent everywhere

TF-IDF

$$tf(w, d) = \log(1 + f(w, d))$$

$$idf(w, D) = \log\left(\frac{N}{f(w, D)}\right)$$

w is a word, d is a document, D is the corpus, $N = |D|$

Intuitions:

- frequent in a single context is good
- avoid infinities
- appearing in every document is bad
- score of 100 (vs. 1) is not 100 times more relevant

TF-IDF

Find tf-idf representation for d1 for the following:

d1 = any big cat

d2 = big cat

d3 = cat dog cat

N = 3 Vocabulary = {any, big, cat, dog}

$$tf(\text{any}, d1) = \log(1+1) = 0.693$$

$$idf(\text{any}, D) = \log(3/1) = 1.099$$

$$tf(\text{big}, d1) = \log(1+1) = 0.693$$

$$idf(\text{big}, D) = \log(3/2) = 0.405$$

$$tf(\text{cat}, d1) = \log(1+1) = 0.693$$

$$idf(\text{cat}, D) = \log(3/3) = 0$$

$$tf(\text{dog}, d1) = \log(1+0) = 0$$

$$idf(\text{dog}, D) = \log(3/1) = 1.099$$

$$tf-idf(\text{any}, d1, D) = 0.693 * 1.099 = 0.761$$

$$tf-idf(\text{big}, d1, D) = 0.693 * 0.405 = 0.281$$

$$tf-idf(\text{cat}, d1, D) = 0.693 * 0 = 0$$

$$tf-idf(\text{dog}, d1, D) = 0 * 1.099 = 0$$

tf-idf representation for d1 is = [0.761 , 0.281 , 0 , 0]

$$tf(w, d) = \log(1 + f(w, d))$$

$$idf(w, D) = \log\left(\frac{N}{f(w, D)}\right)$$

TF-IDF

Find tf-idf representation for d1 for the following:

d1 = any big cat

d2 = big cat

d3 = cat dog cat

N = 3 Vocabulary = {any, big, cat, dog}

Do the same for d2

$$tf(w, d) = \log(1 + f(w, d))$$

$$idf(w, D) = \log\left(\frac{N}{f(w, D)}\right)$$

TF-IDF

Find tf-idf representation for d1 for the following:

d1 = any big cat

d2 = big cat

d3 = cat dog cat

N = 3 Vocabulary = {any, big, cat, dog}

$$\text{tf}(\text{any}, d2) = \log(1+0) = 0$$

$$\text{tf}(\text{big}, d2) = \log(1+1) = 0.693$$

$$\text{tf}(\text{cat}, d2) = \log(1+1) = 0.693$$

$$\text{tf}(\text{dog}, d2) = \log(1+0) = 0$$

$$\text{idf}(\text{any}, D) = \log(3/1) = 1.099$$

$$\text{idf}(\text{big}, D) = \log(3/2) = 0.405$$

$$\text{idf}(\text{cat}, D) = \log(3/3) = 0$$

$$\text{idf}(\text{dog}, D) = \log(3/1) = 1.099$$

$$\text{tf-idf}(\text{any}, d2, D) = 0 * 1.099 = 0$$

$$\text{tf-idf}(\text{big}, d2, D) = 0.693 * 0.405 = 0.281$$

$$\text{tf-idf}(\text{cat}, d2, D) = 0.693 * 0 = 0$$

$$\text{tf-idf}(\text{dog}, d2, D) = 0 * 1.099 = 0$$

tf-idf representation for d2 is = [0 , 0.281 , 0 , 0]

$$\text{tf}(w, d) = \log(1 + f(w, d))$$

$$\text{idf}(w, D) = \log\left(\frac{N}{f(w, D)}\right)$$

IDFs are Same as before,
no need to calculate again

Perfect word representations

- shared lemmas: mouse/mice, dormir/duermes, etc.
- different word senses: computer mouse vs. pet mouse, river bank vs. financial bank, etc.
- synonyms: couch/sofa, car/automobile, etc.
- antonyms: long/short, dark/light, etc.
- word similarity: dog/cat, doctor/nurse, etc.
- word relatedness: cup/coffee, scalpel/surgeon, etc.
- word valence: excited and relaxed are high valence, depressed and angry are low valence
- word arousal: excited and angry are high arousal, relaxed and depressed are low arousal

Neighboring words hint at semantics

- Imagine you didn't know what ignite meant:
 - . . . fusion fire does not ignite till temperatures . . .
 - . . . plumes of flame ignite from the smokestacks . . .
 - . . . over low heat. Ignite with a match ...
- But you had seen another word in similar contexts:
 - . . . the way the fire is lit or the heat source . . .
 - ... flame couldn't have lit a cigarette . . .
 - . . . kiln-dried logs that lit with a match ...

Intuition: if two words are semantically similar, they will appear in text with similar surrounding words

Term-term matrix

A term-term co-occurrence matrix X is a $|V| \times |V|$ matrix where:

- $|V|$ is the number of words in the vocabulary
- each cell $X_{i,j}$ records how often word j occurred in the context of word i
- each row X_i is the vector representation for word i

Context may be defined in different ways:

- The same document
- The same sentence
- Within $\pm n$ words of each other

V is typically the 10,000 - 50,000 most frequent words

- Each word is represented by a large vector

Build a term-term matrix

docs = ["any big cat", "big cat", "**cat** dog **cat**"]

- Context: Within ± 1 words of each other

	any	big	cat	dog
any	0	1	0	0
big	1	0	2	0
cat	0	2	0	2
dog	0	0	2	0

- Context: Within ± 2 words of each other

	any	big	cat	dog
any	0	1	1	0
big	1	0	2	0
cat	1	2	2	2
dog	0	0	2	0

Representing a sentence from term term matrix

docs = ["any big cat", "big cat", "**cat** dog **cat**"]

- Context: Within ± 1 words of each other

	any	big	cat	dog
any	0	1	0	0
big	1	0	2	0
cat	0	2	0	2
dog	0	0	2	0

- “any big cat” will be represented as:
 - Here each word is a vector
 - The whole sentence is represented by the matrix

any	big	cat
0	1	0
1	0	2
0	2	0
0	0	2

Easy way to build a term-term matrix

- Build a binary BoW for the sentences
- Transpose it
- Multiply it with the original matrix
- Voila
- Try it: docs = ["any big cat", "big cat", "cat dog cat"]

Easy way to build a term-term matrix

- Build a binary BoW for the sentences
- Transpose it
- Multiply it with the original matrix

docs = ["any big cat", "big cat", "cat dog cat"]

term	any	big	cat	dog
doc1	1	1	1	0
doc2	0	1	1	0
doc3	0	0	1	1

$M =$

1	1	1	0
0	1	1	0
0	0	1	1

$M^T =$

1	0	0
1	1	0
1	1	1
0	0	1

$M^T . M =$

1	1	1	0
1	2	2	0
1	2	3	1
0	0	1	1

Easy way to build a term-term matrix

docs = ["any big cat", "big cat", "cat dog cat"]

$M^T \cdot M =$

	any	big	cat	dog
any	1	1	1	0
big	1	2	2	0
cat	1	2	3	1
dog	0	0	1	1

- Build a binary BoW for the sentences
- Transpose it
- Multiply it with the original matrix

Note: Here context is same document &

Diagonal values of this matrix represents how many document had that specific word.

Comparing word vectors

- How do you know the vectors you built make any sense?
- You need to compare these vectors
- What techniques do we have?

Cosine similarity

- Most common similarity measure

$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v}^\top \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} = \frac{\overbrace{\sum_i v_i w_i}^{\text{dot product}}}{\underbrace{\sqrt{\sum_i v_i^2}}_{\text{length of } \mathbf{v}} \underbrace{\sqrt{\sum_i w_i^2}}_{\text{length of } \mathbf{w}}}$$

Cosine similarity: Why?

- Range is between 1 and -1
- Why not Euclidean distance?

$$\sqrt{\sum_{i=1}^n (v_i - w_i)^2}$$

$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v}^\top \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} = \frac{\overbrace{\sum_i v_i w_i}^{\text{dot product}}}{\underbrace{\sqrt{\sum_i v_i^2}}_{\text{length of } \mathbf{v}} \underbrace{\sqrt{\sum_i w_i^2}}_{\text{length of } \mathbf{w}}}$$

Let's try this for these three vectors: $\mathbf{u} = [0, 1, 0, 1]$ $\mathbf{v} = [1, 0, 1, 0]$ $\mathbf{w} = [3, 0, 3, 0]$

What is the cosine similarity? What is the Euclidean distance? Which one makes more sense?

What's wrong with term term matrix?

- Sparse. Very sparse.
- Does not carry any contextual information
- Does not represent how important a word is in a sentence

Using word vectors

For word tasks:

- finding synonyms via cosine
- as classifier features when the input is one word

For sentence/document tasks:

- First, combine all word vectors
- You can combine yourself (using centroid technique): usually needed for classical ML models; or
- You can let an RNN handle things
- These vectors can then be used for classification

Sparse vs. dense vectors

Vectors we studied are very sparse

Advantages of small, dense word vectors:

- fewer feature weights to learn in machine learning
- fewer features can reduce overfitting
- forces sharing; there are not enough dimensions for each word to be completely independent

Word embeddings

- Rather than count co-occurrence, let's try to predict it
- We can do it in two ways
 - Predict the target word given the neighboring words: CBOW



- Predict the neighboring words given the target word: Skip-gram



- CBOW is easy, but...
- We will focus on Skip-gram

Skip-gram embeddings

- Input: a word, taken from some text
- Output: the 5 preceding and 5 following words
- Try it!
- Input: hippopotamus
- Output: [?, ?, ?, ?, ?, hippopotamus, ?, ?, ?, ?, ?]

Skip-gram embeddings

- Input: a word, taken from some text
- Output: the 5 preceding and 5 following words
- Try it!
- Input: hippopotamus
- Output: [?, ?, ?, ?, ?, hippopotamus, ?, ?, ?, ?, ?]
- This task is impossible! But that's okay because
 - Creating training data is easy
 - We'll only use word vectors learned as part of training

Creating Training Data is Easy

thou	shalt	not	make	a	machine	in	the	...
------	-------	-----	------	---	---------	----	-----	-----

thou	shalt	not	make	a	machine	in	the	...
------	-------	-----	------	---	---------	----	-----	-----


thou	shalt	not	make	a	machine	in	the	...
------	-------	-----	------	---	---------	----	-----	-----

thou	shalt	not	make	a	machine	in	the	...
------	-------	-----	------	---	---------	----	-----	-----

thou	shalt	not	make	a	machine	in	the	...
------	-------	-----	------	---	---------	----	-----	-----

input word	target word
not	thou
not	shalt
not	make
not	a
make	shalt
make	not
make	a
make	machine
a	not
a	make
a	machine
a	in
machine	make
machine	a
machine	in
machine	the
in	a
in	machine
in	the
in	likeness

Training an embedding model

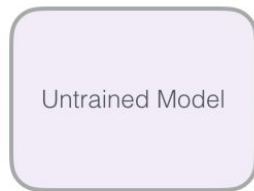


Training an embedding model

Actual
Target

0
0
0
...
0
1
...
0

not



Model
Prediction

0	aardvark	=	0
0	aarhus		0
0.001	aaron		-0.001
...			...
0.4	taco		-0.4
0.001	thou		0.999
...			...
0.0001	zyzzyva		-0.0001

Error


Update
Model
Parameters



Problem with this model

- Step 3 is very expensive as we have to project the output into the entire vocabulary space
 - Especially, we have to do it for every single step

Let's switch the task: Mikolov's entry



This converts a complex neural learning task into a simple log-linear binary classification task

Converting the data for this task

input word	target word
not	thou
not	shalt
not	make
not	a
make	shalt
make	not
make	a
make	machine

input word	output word	target
not	thou	1
not	shalt	1
not	make	1
not	a	1
make	shalt	1
make	not	1
make	a	1
make	machine	1

Converting the data for this task

input word	target word
not	thou
not	shalt
not	make
not	a
make	shalt
make	not
make	a
make	machine

All one's is
no good for
learning

input word	output word	target
not	thou	1
not	shalt	1
not	make	1
not	a	1
make	shalt	1
make	not	1
make	a	1
make	machine	1

Negative sampling

Skipgram


shalt	not	make	a	machine
input		output		
make		shalt		
make		not		
make		a		
make		machine		

Negative Sampling

input word	output word	target
make	shalt	1
make	aaron	0
make	taco	0

Word2vec Training

- we determine the size of our vocabulary
- create two matrices – an Embedding matrix and a Context matrix
- Initialize these matrices with random values



Word2vec Training

input word	output word	target
not	thou	1
not	aaron	0
not	taco	0
not	shalt	1
not	mango	0
not	finglonger	0
not	make	1
not	plumbus	0
...	...	...


Embedding

				aardvark
				aarhus
				aaron
				...
				not
				...
				...
				...
				zyzzyva

Context

				aardvark
				aarhus
				aaron
				...
				taco
				...
				thou
				...
				zyzzyva



Look up
embeddings




not

aaron
taco
thou

Word2vec Training

input word	output word	target	input • output	sigmoid()	Error
not 	thou 	1	0.2	0.55	0.45
not 	aaron 	0	-1.11	0.25	-0.25
not 	taco 	0	0.74	0.68	-0.68



Update
Model
Parameters

Using embeddings

It's rarely necessary to train skip-gram or GloVe directly.

Download pre-trained word embeddings:

- Skip-gram <https://code.google.com/archive/p/word2vec/>
- GloVe <https://nlp.stanford.edu/projects/glove/>

Many models provide pre-trained embeddings:


<https://github.com/Hironsan/awesome-embedding-models>

Which one should I choose? Try a few and see what works!

Semantic properties of embeddings

- Different types of similarity or association
 - Based on the context window, word association changes
 - smaller window sizes (2-15) lead to embeddings where high similarity scores between two embeddings indicates that the words are interchangeable
 - Larger window sizes (15-50, or even more) lead to embeddings where similarity is more indicative of relatedness of the words
- Analogy/relational similarity
 - Parallelogram model: Apple is to Tree as Grape is to _____
- Historical context

Historical semantic context



Inspecting embeddings

How do I know if my word embeddings make sense?

- Check by hands
- Project the words to a visible dimension
- Use linear algebra

Visualizing embeddings

We usually have very high-dimensional vectors for each words. t-SNE can project down to 2.



Algebra

```
>>> cosine(vector("queen"), vector("king"))
0.7252606
>>> cosine(vector("queen"),
... vector("king")-vector("man")+vector("woman"))
0.7880841
>>> cosine(vector("Paris"), vector("Rome"))
0.58241177
>>> cosine(vector("Paris"),
... vector("Rome")-vector("Italy")+vector("France"))
0.71733016
```

Other standard evaluations

Correlation with human judgments of similarity

- **WordSim-353** noun similarity, e.g., (plane, car, 5.77)
- **SimLex-999** adjective, noun, and verb similarities
- **SCWS** word similarity given sentential context
- **STS** sentence-level similarity

Accuracy at similarity-based task

- **TOEFL** e.g., Levied is closest in meaning to: imposed, believed, requested, correlated
- **analogies** e.g., Athens is to Greece as Oslo is to _____

Bias in embeddings

Embeddings reflect the language they were trained on

```
>>> cosine(vector("attractive"), vector("man"))
0.3085765
>>> cosine(vector("attractive"), vector("woman"))
0.41110972
>>> cosine(vector("dumb"), vector("American"))
0.41180187
>>> cosine(vector("dumb"), vector("European"))
0.26587355
```

Contextual word embeddings

Traditional word vectors ignore context

The river bank: [0.3, -0.1, -0.2] [0.1, -0.3, -0.2] [-0.6, 0.3, -0.1]

A bank deposit: [0.0 , 0.0 , -0.2] [-0.6, 0.3, -0.1] [-0.3, -0.3, 0.0]


Should these two banks really have the same vectors?

Contextual word embeddings

Word embeddings

Input 1 word


Output 1 embedding



Contextual word embeddings

Input n words

Output n embeddings



Learning contextual word embeddings

We need to make up a prediction task that

- takes n words as input
- produces n vectors as output
- requires only unlabeled data

ELMo's task: language modeling


What is a language model?

- Given a sequence of words, what is the next most probable word?
- Unsupervised, great for learning representations

ELMo combines a forward language model and a backward language model.

Transformers use the same idea, but in a much larger canvas.

ELMo's task: language modeling



How to use contextual word embeddings?

Contextual word embeddings are trained on unlabeled data. How do we use them on the task we care about?

- Extract word vectors, use as features
- Fine-tune contextual embedding model, i.e., continue training the model, but now on our labeled data instead of the unlabeled data